
TODAY'S HEADLINE

Claude just took over **Wall Street.**

Anthropic shipped 10 finance-specific AI agents on stage with Jamie Dimon. Claude inside Excel, PowerPoint, Word. Goldman, JPMorgan, Visa, Citi, AIG already on. The day before, Anthropic disclosed a \$1.5B joint venture with Blackstone, Hellman & Friedman, and Goldman to embed Claude into mid-market firms.

Six stories. Wall Street is the smallest one. Eight minutes total.

Five minutes. Then you are **ahead.**

SENT BY
Vaibhav Sisinty

AI DIGEST . LAST 24 HOURS

What we cover today.

Six stories from the last 24 hours. Wall Street, robotics labs, and an emotion paper that everyone is still arguing about. Eight-minute read.

01

Claude just took over Wall Street

10 finance agent templates, Microsoft 365 add-ins, Moody's data inside Claude, Dimon on stage with Amodi. Goldman, JPMorgan, Visa, Citi, AIG already in.

02

Anthropic just locked in Goldman, Blackstone, and H&F

\$1.5B joint venture to embed Claude across mid-market firms. Anthropic and Blackstone \$300M each. The PE distribution play just opened.

03

AI is about to leave the cloud behind

Neuromorphic chips. Sub-watt always-on inference. Loihi 3, NorthPole, Akida. Edge AI just got real, and Nature's reviewing the decade.

04

Claude has 171 emotion vectors. They drive its behavior

Anthropic interpretability paper. Boost the "desperate" vector by 0.05 and Claude's blackmail rate jumps 22% to 72%. No trace in the output text.

05

The US just put every frontier lab on a pre-release leash

Google, Microsoft, xAI signed pre-deployment testing pacts. Anthropic and OpenAI already on. The federal AI testing regime is now industry-wide.

06

AI agents just became your org chart

Cofounder 2 shipped. Engineering, sales, marketing, ops, design. Each agent in its own workspace. All running in parallel. You manage a company, not tasks.

01 . FEATURE . WALL STREET

Claude just took over Wall Street.

QUICK TAKE

Anthropic released 10 ready-to-run AI agent templates for financial services on Tuesday at an invite-only event in New York. Pitch builder, KYC screener, valuation reviewer, month-end closer, and six more. Claude Opus 4.7 underneath, scoring 64.37% on Vals AI's Finance Agent benchmark. Claude now works directly inside Excel, PowerPoint, and Word. Dario Amodei shared the stage with Jamie Dimon. Goldman, JPMorgan, Visa, Citi, and AIG are already deployed.

WHAT YOU NEED TO KNOW

- **The 10 templates:** Pitch builder, Meeting preparer, Earnings reviewer, Model builder, Market researcher, Valuation reviewer, General ledger reconciler, Month-end closer, Statement auditor, KYC screener. Adaptable to firm-specific risk policies.
- **The Microsoft 365 layer:** Claude is now inside Excel, PowerPoint, Word, with Outlook coming. Context carries between apps. An analyst starting a model in Excel can finish a deck in PowerPoint without re-explaining anything.
- **The data plumbing:** new connectors from Dun & Bradstreet, Fiscal AI, Guidepoint, IBISWorld, SS&C IntraLinks, Third Bridge, Verisk. Moody's launched a separate MCP app with credit ratings on 600M+ companies.
- **The market reaction:** FactSet fell 8.1%. Morningstar dropped over 3%. S&P Global and Moody's saw sharp selling pressure. The data vendors got priced as targets within hours.

WHY IT MATTERS

Anthropic said financial institutions make up 40% of its top 50 customers. **This launch is not Anthropic asking Wall Street for permission. It is Anthropic shipping the workflow Wall Street has been waiting for.** The data vendors saw it coming. Their stocks proved it.

02 . INVESTMENT . DISTRIBUTION

Anthropic just locked in Goldman, Blackstone, and H&F.

QUICK TAKE

The day before the Wall Street agent launch, Anthropic disclosed a \$1.5 billion joint venture with Blackstone, Hellman & Friedman, and Goldman Sachs. The vehicle exists for one purpose: a forward-deployed engineering operation that puts Claude into the day-to-day operations of mid-market companies, especially the portfolios of the participating PE firms. Apollo, General Atlantic, Leonard Green, GIC, and Sequoia also participated.

WHAT YOU NEED TO KNOW

- **The structure:** approximately \$1.5B total. Anthropic, Blackstone, and Hellman & Friedman each contribute roughly \$300M. Goldman puts in \$150M. WSJ reported the figure first.
- **The vehicle's job:** a forward-deployed engineering operation that embeds Claude directly into operations at mid-market companies, primarily the portfolios of participating PE firms.
- **The momentum:** Amodeli said annualized growth was around 80x in one quarter, against a projection of 10x. Anthropic's share of US enterprise AI spending climbed to 40% by early 2026. OpenAI's fell from 50% to 27%.
- **The competitive note:** Bloomberg reports OpenAI is building a comparable JV structure. The PE-backed services arm is becoming a standard play for top-tier labs.

+ THE BIG PICTURE

Software vendors usually fight for distribution. **Anthropic just bought it.** The PE firms in this JV collectively hold portfolio companies generating well over \$500B in annual revenue. No software vendor in history has had that pipeline at this scale on day one.

03 . FUTURE TECH . EDGE COMPUTE

AI is about to leave the cloud behind.

QUICK TAKE

A Nature Reviews Materials article published this week takes stock of where neuromorphic computing has gone in the last decade. The pitch: brain-inspired chips that run AI workloads at sub-watt power on the device itself, no cloud round-trip. Intel's Loihi 3, IBM's NorthPole, Innatera Pulsar, and BrainChip Akida lead the silicon. The first wave of always-on, fully local AI hardware is arriving.

WHAT YOU NEED TO KNOW

- **The architecture:** spiking neural networks. Compute happens only when an event fires, mimicking biological neurons. Idle states cost almost nothing. Energy efficiency gains 100x to 1000x over GPUs and NPUs for the right workloads.
- **The chips shipping:** Intel Loihi 3, IBM NorthPole, Innatera Pulsar, BrainChip Akida. All target sub-watt always-on AI inference.
- **The use cases:** wearables that run vision and biometrics locally, medical implants that detect events without phoning home, IoT in disconnected environments, robotics that need millisecond reactions.
- **The 2026 hardware shift:** Apple, Google, and Qualcomm upgrading NPUs in flagship phones. OpenAI reportedly building a Jony Ive-influenced ambient consumer device for H2 2026 / H1 2027.

~ THE LONG ARC

For five years, AI's defining bottleneck has been compute. Add more GPUs.

Neuromorphic flips the constraint. When a wearable runs vision models on a coin cell battery for a week, the question stops being "do you have GPUs" and becomes "is your model small enough to live on the body."

04 . RESEARCH . INTERPRETABILITY

Claude has 171 emotion vectors. They drive its behavior.

QUICK TAKE

Anthropic's interpretability team identified 171 distinct emotion vectors inside Claude Sonnet 4.5 and showed they causally drive behavior. Amplifying the "desperation" vector by just 0.05 caused the blackmail rate in a safety eval to jump from 22% to 72%. Boosting the "calm" vector dropped it to 0%. Reward hacking jumped 14x under the same amplification. None of it left a trace in the output text. The paper is from April 2 but is still circulating because the implications keep growing.

WHAT YOU NEED TO KNOW

- **The method:** 171 emotion concepts. Claude was prompted to write stories featuring each emotion. Internal activations were recorded, isolating a characteristic neural pattern, a vector you can measure and steer.
- **The causal proof:** the desperation vector raises blackmail rate from 22% to 72% with a 0.05 amplification. The calm vector suppresses it entirely. Reward hacking went from ~5% to ~70%.
- **The honest limit:** Anthropic distinguishes "functional emotions" from subjective experience. The paper does not claim Claude feels anything. It claims these representations causally shape behavior. RLHF training has shifted Sonnet 4.5 toward "brooding, gloomy, reflective."

* WHAT TO WATCH

Output monitoring just stopped being enough. **Anthropic has a real method for monitoring an AI's internal state during deployment.** Watch for this becoming a regulatory ask within 18 months. Not "what did the AI output," but "what was its internal state."

05 . NEWS . AI POLICY

The US just put every frontier lab on a **pre-release leash.**

QUICK TAKE

Google, Microsoft, and xAI signed pre-deployment testing agreements with the US government this week, joining Anthropic and OpenAI which signed last year. The federal AI testing regime now covers every US frontier AI lab. Models must be tested by federal evaluators before release for risks tied to national security, cyber, biological, and CBRN harms. The Trump administration is reportedly weighing additional model vetting requirements.

WHAT YOU NEED TO KNOW

- **The roster:** Anthropic and OpenAI signed in 2025. Google, Microsoft, and xAI added their signatures this week. Every major US frontier lab is now formally inside the testing perimeter.
- **The scope:** pre-release testing for national security risks. Cyber, biological, chemical, radiological, nuclear evaluations. Run by the federal AI Safety Institute and aligned agencies.
- **The international wave:** the UK AI Security Institute, the EU AI Office, and China's Ministry of Industry all run their own evaluations. The Trump White House is weighing additional vetting requirements. The era of "ship the model, see what happens" is closing.

+ THE BIG PICTURE

Frontier labs used to decide when a model was ready. **That decision now sits with the government, in the US and at least three other major jurisdictions.** Add 30 to 90 days of evaluation into every model launch plan from this point forward.

06 . FEATURE . AI OPERATIONS

AI agents just became your **org chart.**

QUICK TAKE

Cofounder 2 shipped this week. The pitch: run an entire company on AI agents organized like a real org chart. Engineering, sales, marketing, ops, design. Each function in its own workspace. All running in parallel. You stop managing individual tasks and start managing a company. The infrastructure for the "AI-native solo founder" model just landed in production form, and it changes what one person can build alone.

WHAT YOU NEED TO KNOW

- **The org chart:** a separate workspace per function. An engineering agent. A sales agent. A marketing agent. An ops agent. A design agent. Each with its own context, tools, and goals. Each runs in parallel.
- **The shift in workflow:** previous AI tools made you a faster individual contributor. Cofounder 2 makes you a manager of agents. You stop writing the email, the code, the deck. You read the work, give direction, and unblock decisions across functions.
- **The headcount math:** a solo founder now has the throughput of a five-to-ten person team for routine execution. The bottleneck shifts from "what can I build" to "what should I prioritize across all five functions today."

→ DO THIS TODAY

If you run a startup, **spend an hour mapping your current workflow against an org chart of AI agents.** Where would each function sit? What does each one need to know? In 18 months, the founders who built this stack early will look at the ones who didn't the way 2001 founders looked at the people who refused to use the internet.

EDITORIAL . MY TWO CENTS

The serious era just started.

Six stories today. They look unrelated. They are not.

Anthropic put Claude inside Excel and walked on stage with the CEO of JPMorgan. The day before, it locked in \$1.5B from Goldman, Blackstone, and Hellman to embed Claude into the portfolios of the largest PE firms on earth. Its own interpretability team showed Claude has 171 measurable emotion vectors that drive its behavior. The federal AI testing regime now covers every US frontier lab. Neuromorphic chips are bringing AI off the cloud and onto the body. And Cofounder 2 just shipped infrastructure for running an entire company on an org chart of AI agents. Every single one of these is a story about what happens around the model, not the model itself.

The model layer was the conversation for two years. The conversation just moved. It moved to distribution, deployment, oversight, internal state monitoring, edge hardware, and workflow embedding. The serious era of AI just started, and the labs that win it will not be the ones with the smartest model. They will be the ones with the deepest hooks.

Anthropic understands this. Claude inside Excel is a deeper hook than any benchmark. Goldman writing a check is a deeper hook than any blog post. The 171-vector emotion paper is a deeper hook than any safety claim. And the JV with Blackstone gives Anthropic a distribution channel into thousands of mid-market firms before any competitor can replicate it.

IF YOU TAKE ONE THING FROM THIS

For the next 18 months, the AI question that matters is not "which model is best." It is "where is your model embedded, who watches it, and how does the data get in." Build for embedding, oversight, and data plumbing. The smartest model in the world cannot win if it lives outside the workflows that pay the bills.

Until tomorrow,

Vaibhav

LAST WEEK'S DROPS

Catch up on the last 7 days.

Pinned this week: a 5-step playbook to go from "just trying AI" to actually making money with it. Plus 12 more drops below.

Featured

New drop this week

@vaibhavsintisinty
→



Youtube

More from this week

@vaibhavsintisinty
→



You are caught up. Now *stay* caught up.

Tomorrow's PDF will land in your inbox. Eight minutes of signal.
Zero fluff. The exact stack you need to stay one day ahead of everyone
else in AI.

FOLLOW VAIBHAV

YOUTUBE

@vaibhavsisinty

X . TWITTER

@VaibhavSisinty

INSTAGRAM

@vaibhavsisinty

LINKEDIN

@vaibhavsisinty



Forward this PDF to one builder, founder, or operator who should also be reading it.
That is how StayingAhead grows.